

Handling deadlocks

- definition, wait-for graphs
- fundamental causes of deadlocks
- resource allocation graphs and conditions for deadlock existence
- approaches to handling deadlocks
 - ♦ deadlock prevention
 - ♦ deadlock avoidance
 - ♦ deadlock detection/resolution
 - centralized algorithms
 - distributed algorithms
 - hierarchical algorithms

1

What is a deadlock?

- deadlock – a set of processes is blocked waiting for requirements that cannot be satisfied
- illustrated by a wait-for-graph (WFG)
 - ♦ nodes – processes in the system
 - ♦ directed edges – wait-for blocking relation
 - ♦ a cycle in WFG indicates a deadlock
- starvation – a process is indefinitely prevented from making progress
 - ♦ deadlock implies starvation, is the converse true?

2

Fundamental causes of deadlocks

- Mutual exclusion** — if one process holds a resource, other processes requesting that resource must wait until the process releases it (only one can use it at a time)
- Hold and wait** — processes are allowed to *hold one* (or more) resource and be *waiting* to acquire additional resources that are being held by other processes
- No preemption** — resources are released voluntarily; neither another process nor the OS can force a process to release a resource
- Circular wait** — there must exist a set of waiting processes such that P₀ is waiting for a resource held by P₁, P₁ is waiting for a resource held by P₂, ... P_{n-1} is waiting for a resource held by P_n, and P_n is waiting for a resource held P₀


3

Resource allocation graph

- The deadlock conditions can be modeled using a directed graph called a *resource-allocation graph* (RAG)
 - ♦ 2 kinds of nodes:
 - *Boxes* — represent resources
 - Instances of the resource are represented as dots within the box
 - *Circles* — represent processes
 - ♦ 2 kinds of (directed) edges:
 - *Request edge* — from thread to resource — indicates the thread has requested the resource, and is waiting to acquire it
 - *Assignment edge* — from resource instance to thread — indicates the thread is holding the resource instance
 - ♦ When a request is made, a request edge is added
 - When request is fulfilled, the request edge is transformed into an assignment edge
 - When process releases the resource, the assignment edge is deleted

4


RAG with single resource instances



- a cycle in RAG with single resource instances is necessary and sufficient for deadlock

5

RAG with multiple resource instances



- cycle does not indicate deadlock
- knot – strongly connected subgraph (no sinks) with no outgoing edges
- a knot in RAG is necessary and sufficient for deadlock

6

Deadlock prevention and avoidance

- **Deadlock prevention** — eliminate one of the 4 deadlock conditions
 - ◆ examples
 - acquire all resources before proceeding (no wait while hold)
 - allow preemption (eliminate 3d condition)
 - prioritize processes and assign resources in the order of priorities (no circular wait)
 - ◆ may be inefficient
- **Deadlock avoidance** — consider each resource request, and only fulfill those that will not lead to deadlock
 - ◆ Stay in a *safe state* — a state with no deadlock where resource requests can be granted in some order such that all processes will complete
 - ◆ may be inefficient
 - Must know resource requirements of all processes in advance
 - Resource request set is known and fixed, resources are known and fixed
 - Complex analysis for every request

7

Deadlock detection

- **Deadlock detection and resolution** — detect, then break the deadlock
- detection
 - ◆ issues
 - maintenance of WFG
 - search of WFG for deadlocks
 - ◆ requirements
 - progress — no undetected deadlocks
 - safety — no false (phantom) deadlocks
- resolution
 - ◆ roll back one or more processes to break dependencies in WFG and resolve deadlocks

8

Distributed deadlock detection algorithms

- Centralized algorithm - coordinator maintains global WFG and searches it for cycles
 - ◆ simple algorithm
 - ◆ Ho and Ramamoorthy's one- and two-phase algorithms
- Distributed algorithms - Global WFG, with responsibility for detection spread over many sites
 - ◆ Obermarck's path-pushing
 - ◆ Chandy, Misra, and Haas's edge-chasing
 - ◆ diffusion
- Hierarchical algorithms - hierarchical organization, site detects deadlocks involving only its descendants
 - ◆ Menasce and Muntz's algorithm
 - ◆ Ho and Ramamoorthy's algorithm


9

Simple centralized deadlock detection

- A central coordinator maintains a global wait-for graph (WFG) for the system
- All sites request and release resources (even local resources) by sending request and release messages to the coordinator
- When coordinator receives a request/release, it
 - ◆ updates the global WFG
 - ◆ checks for deadlocks
- problems
 - ◆ large communication overhead, coordinator is a performance bottleneck and single point of failure, etc.
 - ◆ may report phantom deadlock

10

Problem of False Deadlock



- Now assume process p1 releases resource p3 is waiting on
- Slightly thereafter, p2 requests resource p3 is holding
- However, first message reaches coordinator after second message
- The global WFG now has a *false cycle*, which leads to a report of *false deadlock*

11

Ho and Ramamoorthy two phase centralized deadlock detection

- Every site maintains a status table, containing status of all local processes
 - ◆ Resources held, resources waiting on
- Periodically, coordinator requests all status tables, builds a WFG, and searches it for cycles
 - ◆ No cycles - no deadlock
 - ◆ If cycle is found, coordinator again requests all status tables, again builds a WFG, but this time uses only those edges common to both sets of status tables
- Rationale was that by using information from two consecutive reports, coordinator would get a consistent view of the state
 - ◆ However, it was later shown that a deadlock in this WFG does not imply a deadlock exists (see 1 phase alg.)
 - ◆ So, the HR-two-phase algorithm may reduce the possibility of reporting false deadlocks, but doesn't eliminate it

13

Ho and Ramamoorthy one-phase centralized deadlock detection

- Every site maintains two tables
 - ◆ all local processes and resources the locked
 - ◆ resources locked at this site (by both local and non-local processes)
- one site periodically requests both tables (once) constructs WFG; WFG only includes the info on non local processes if this info is matched by the process site and resource site
 - ◆ if cycle – deadlock
 - ◆ if not – no deadlock
- correctly detects deadlocks by eliminating inconsistency of reporting due to message propagation delay
- more space overhead than 2-phase H&R